

Lab04 - (Extra) Plotting with Tensorboard

To observe if the model is training well and identify if the model has converged, it is important to monitor the loss function vs iteration. To do this, we can use [Tensorboard](#) which is a tool designed for visualizing the results of neural network training runs. It provides the visualization needed for checking your machine learning model. Two main uses of tensorboard are to

1. Visualize the constructed computational graph via `writer.add_graph` method.
2. Track and visualize metrics such as loss and accuracy via `writer.add_scaler` method

Objectives:

1. Learn how to use Tensorboard to monitor training metrics such as training loss and accuracy and visualize the graph for our model.

Task:

Visualize the training of the CIFAR10 model using Tensorboard.

Content:

1. [Load Training Set and Construct Network](#)
2. [Visualize a Graph](#)
3. [Monitoring a Scalar Value](#)
4. [Comparing a Scalar Across Different Runs](#)
5. [Grouping Plots](#)
6. [Plotting Multiple Plots in the Same Figure](#)

```
In [1]: from google.colab import drive  
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
In [2]: cd "/content/gdrive/My Drive/UCCD3074_Labs/UCCD3074_Lab4"
```

/content/gdrive/My Drive/UCCD3074_Labs/UCCD3074_Lab4

```
In [3]: import numpy as np
import matplotlib.pyplot as plt

from cifar10 import CIFAR10

import torch
import torch.nn as nn
import torch.nn.functional as F

import torch.optim as optim

import torchvision
import torchvision.transforms as transforms

from torch.utils.data import DataLoader

import time

%load_ext autoreload
%autoreload 2
```

1. Load Training Set and Construct Network

The following code loads the train and test set for CIFAR10.

```
In [4]: transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])
```

```
In [16]: trainset = CIFAR10(train=True, transform=transform, num_samples=10000, download=True)
testset = CIFAR10(train=False, transform=transform, num_samples=2000, download=True)
classes = trainset.classes
```

```
trainloader = DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)
testloader = DataLoader(testset, batch_size=4, shuffle=True, num_workers=2)
```

Files already downloaded and verified

Files already downloaded and verified

The following code constructs the network model

In [6]:

```
class Network(nn.Module):

    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(3*32*32, 50)
        self.relu1 = nn.ReLU()
        self.fc2 = nn.Linear(50, 10)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = self.fc1(x)
        x = self.relu1(x)
        x = self.fc2(x)
        return x

def build_model():
    net = Network()
    if torch.cuda.is_available():
        net = net.cuda()
    return net
```

In [7]:

```
net = build_model()
```

2. Visualize a Graph

One of TensorBoard's strengths is its ability to visualize complex model structures. Let's visualize the model we built. First, we import `SummaryWriter` from `torch.utils.tensorboard` and define a [SummaryWriter](#) object:

```
writer = SummaryWriter(log_dir)
```

where `log_dir` is the directory to store the written item.

```
In [40]: from torch.utils.tensorboard import SummaryWriter

writer = SummaryWriter('runs/graph')
```

To create the computational graph, `summarywriter` requires you to supply a batch data.

```
In [41]: X = torch.randn(4, 3, 32, 32)
if torch.cuda.is_available():
    X = X.cuda()
```

Now, let's add the graph to tensorboard through the command `add_graph`

```
In [42]: writer.add_graph(net, X)
```

After we have finish updating the graph, close the writer.

```
In [43]: writer.close()
```

Now, let's open up tensorboard on Colab to view the graph that we have just writtten to `runs/graph` .

```
In [56]: %load_ext tensorboard
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

```
In [76]: %tensorboard --logdir runs
```

Output hidden; open in <https://colab.research.google.com> to view.

3. Monitoring a Scalar Value

Besides displaying a graph, Tensorboard can be used to monitor a scalar value over different time step. We can use it to monitor the training/testing accuracy/loss.

Add scalar during training

To monitor a variable, call the function `add_scalar` to update the variable continuously during training.

```
writer.add_scalar(tag, scalar_value, step)
```

where the input arguments are:

- `tag` (string) is the name of the experiment. Tensorboard can show multiple logs from *different* experiments with the *same* tag on the same figure. This allows easy comparison.
- `scalar_value` (a scalar value) is the value to be saved and monitored, e.g., the accuracy or loss value.
- `step` is the training step when `scalar_value` is generated.

In the `train` function below, we monitor the training loss as the training commences. We shall save the `train_loss` to `tensorboard` (line 54) every 100 batch cycles (line 47).

```
In [59]: import time

def train (model, writer, lr=0.001, num_epochs=5, loop_per_val=200):

    # set the optimizer
    optimizer = optim.SGD(net.parameters(), lr=lr)

    # set the criterion
    criterion = nn.CrossEntropyLoss()

    # set to training mode
    net.train()

    # train the model
    for e in range(num_epochs):

        running_loss = 0
        running_count = 0
```

```

for i, (inputs, labels) in enumerate(trainloader):

    # transfer input to GPU
    if torch.cuda.is_available():
        inputs = inputs.cuda()
        labels = labels.cuda()

    # set grad to zero
    net.zero_grad()

    # forward propagation
    outputs = net(inputs)

    # compute loss
    loss = criterion(outputs, labels)

    # backward propagation
    loss.backward()

    # update model
    optimizer.step()

    # compute running loss and validation
    running_loss += loss.item()
    running_count += 1

    # display the averaged loss value every 100 cycles
    if i % loop_per_val == loop_per_val - 1:
        train_loss = running_loss / loop_per_val
        running_loss = 0.
        running_count = 0
        print(f'[Epoch {e+1:2d}/{num_epochs:d} Iter {e*len(trainloader)+i:5d}]: train_loss = {train_loss:.4f}')

    # Update tensorboard
    writer.add_scalar("Training loss", train_loss, e*len(trainloader)+i)

return model

```

Next, create a summary writer object. We shall name our `log_dir` based on the learning rate (`lr`) settings since we will be comparing different `lr` settings later.

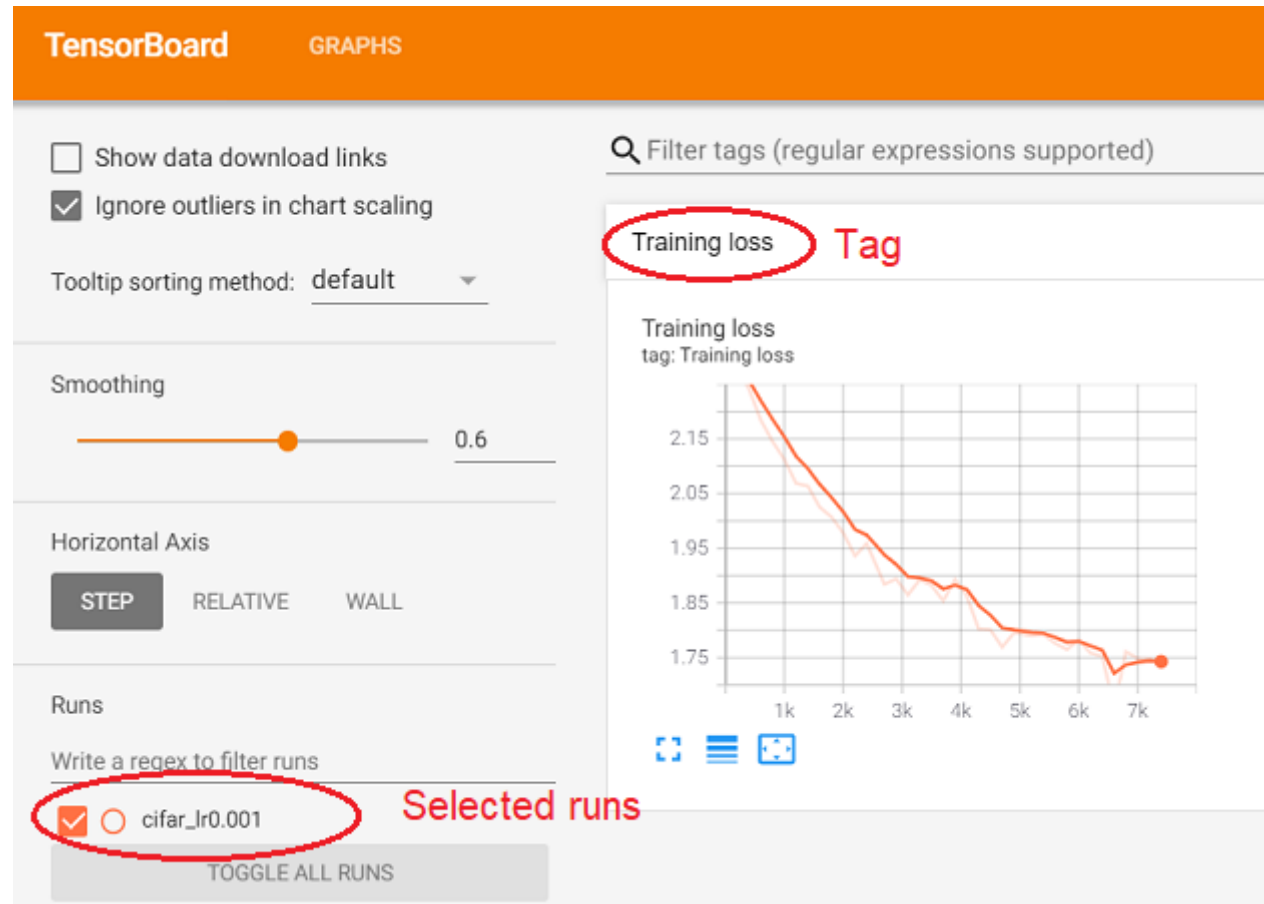
```
In [60]: writer = SummaryWriter('runs/cifar_lr0.001')
```

```
In [61]: net = build_model()  
net = train (net, writer, lr = 0.001, num_epochs=3, loop_per_val=200)
```

```
[Epoch 1/3 Iter 199]: train_loss = 2.2629  
[Epoch 1/3 Iter 399]: train_loss = 2.2043  
[Epoch 1/3 Iter 599]: train_loss = 2.1585  
[Epoch 1/3 Iter 799]: train_loss = 2.1023  
[Epoch 1/3 Iter 999]: train_loss = 2.0967  
[Epoch 1/3 Iter 1199]: train_loss = 2.0120  
[Epoch 1/3 Iter 1399]: train_loss = 2.0342  
[Epoch 1/3 Iter 1599]: train_loss = 2.0096  
[Epoch 1/3 Iter 1799]: train_loss = 1.9712  
[Epoch 1/3 Iter 1999]: train_loss = 1.9607  
[Epoch 1/3 Iter 2199]: train_loss = 1.9945  
[Epoch 1/3 Iter 2399]: train_loss = 1.9643  
[Epoch 2/3 Iter 2699]: train_loss = 1.8742  
[Epoch 2/3 Iter 2899]: train_loss = 1.8986  
[Epoch 2/3 Iter 3099]: train_loss = 1.8525  
[Epoch 2/3 Iter 3299]: train_loss = 1.9143  
[Epoch 2/3 Iter 3499]: train_loss = 1.8484  
[Epoch 2/3 Iter 3699]: train_loss = 1.8375  
[Epoch 2/3 Iter 3899]: train_loss = 1.8288  
[Epoch 2/3 Iter 4099]: train_loss = 1.7753  
[Epoch 2/3 Iter 4299]: train_loss = 1.7768  
[Epoch 2/3 Iter 4499]: train_loss = 1.8118  
[Epoch 2/3 Iter 4699]: train_loss = 1.8389  
[Epoch 2/3 Iter 4899]: train_loss = 1.8338  
[Epoch 3/3 Iter 5199]: train_loss = 1.7953  
[Epoch 3/3 Iter 5399]: train_loss = 1.7889  
[Epoch 3/3 Iter 5599]: train_loss = 1.7601  
[Epoch 3/3 Iter 5799]: train_loss = 1.7343  
[Epoch 3/3 Iter 5999]: train_loss = 1.7597  
[Epoch 3/3 Iter 6199]: train_loss = 1.7446  
[Epoch 3/3 Iter 6399]: train_loss = 1.7255  
[Epoch 3/3 Iter 6599]: train_loss = 1.7353  
[Epoch 3/3 Iter 6799]: train_loss = 1.7630  
[Epoch 3/3 Iter 6999]: train_loss = 1.7387  
[Epoch 3/3 Iter 7199]: train_loss = 1.7192  
[Epoch 3/3 Iter 7399]: train_loss = 1.7349
```

```
In [62]: writer.close()
```

Scroll up to the Tensorboard cell above and check the plot. It should look like the following figure.



4. Comparing a Scalar Across Different Runs

In the following, we shall use a higher learning rate (0.01). Tensorboard will visualize all experiments with the same **tag** in the same plot. This allows you to compare the results from different experiments in the same figure. In this case, the tag name is `Training loss`.

```
In [63]: net = build_model()
writer = SummaryWriter('runs/cifar_lr0.01')
```



```
net = train (net, writer, lr = 0.01, num_epochs=3, loop_per_val=200)
writer.close()
```

```
[Epoch 1/3 Iter 199]: train_loss = 2.1447
[Epoch 1/3 Iter 399]: train_loss = 2.0125
[Epoch 1/3 Iter 599]: train_loss = 1.8855
[Epoch 1/3 Iter 799]: train_loss = 1.8521
[Epoch 1/3 Iter 999]: train_loss = 1.8467
[Epoch 1/3 Iter 1199]: train_loss = 1.8424
[Epoch 1/3 Iter 1399]: train_loss = 1.8056
[Epoch 1/3 Iter 1599]: train_loss = 1.7517
[Epoch 1/3 Iter 1799]: train_loss = 1.7759
[Epoch 1/3 Iter 1999]: train_loss = 1.6565
[Epoch 1/3 Iter 2199]: train_loss = 1.7712
[Epoch 1/3 Iter 2399]: train_loss = 1.7342
[Epoch 2/3 Iter 2699]: train_loss = 1.5594
[Epoch 2/3 Iter 2899]: train_loss = 1.7206
[Epoch 2/3 Iter 3099]: train_loss = 1.6957
[Epoch 2/3 Iter 3299]: train_loss = 1.5883
[Epoch 2/3 Iter 3499]: train_loss = 1.5971
[Epoch 2/3 Iter 3699]: train_loss = 1.6264
[Epoch 2/3 Iter 3899]: train_loss = 1.6380
[Epoch 2/3 Iter 4099]: train_loss = 1.5705
[Epoch 2/3 Iter 4299]: train_loss = 1.5917
[Epoch 2/3 Iter 4499]: train_loss = 1.6448
[Epoch 2/3 Iter 4699]: train_loss = 1.6445
[Epoch 2/3 Iter 4899]: train_loss = 1.6004
[Epoch 3/3 Iter 5199]: train_loss = 1.4905
[Epoch 3/3 Iter 5399]: train_loss = 1.4861
[Epoch 3/3 Iter 5599]: train_loss = 1.5028
[Epoch 3/3 Iter 5799]: train_loss = 1.5535
[Epoch 3/3 Iter 5999]: train_loss = 1.5411
[Epoch 3/3 Iter 6199]: train_loss = 1.4869
[Epoch 3/3 Iter 6399]: train_loss = 1.5098
[Epoch 3/3 Iter 6599]: train_loss = 1.5594
[Epoch 3/3 Iter 6799]: train_loss = 1.5518
[Epoch 3/3 Iter 6999]: train_loss = 1.5819
[Epoch 3/3 Iter 7199]: train_loss = 1.5703
[Epoch 3/3 Iter 7399]: train_loss = 1.5660
```

Scroll up to the Tensorboard cell to observe the result. You should see something similar to the figure below. Tensorboard clearly allows you to compare the monitored value for different runs with different settings easily. To load the data automatically, enable "reload data" and set the reload period to the minimum settings of 30.

Settings

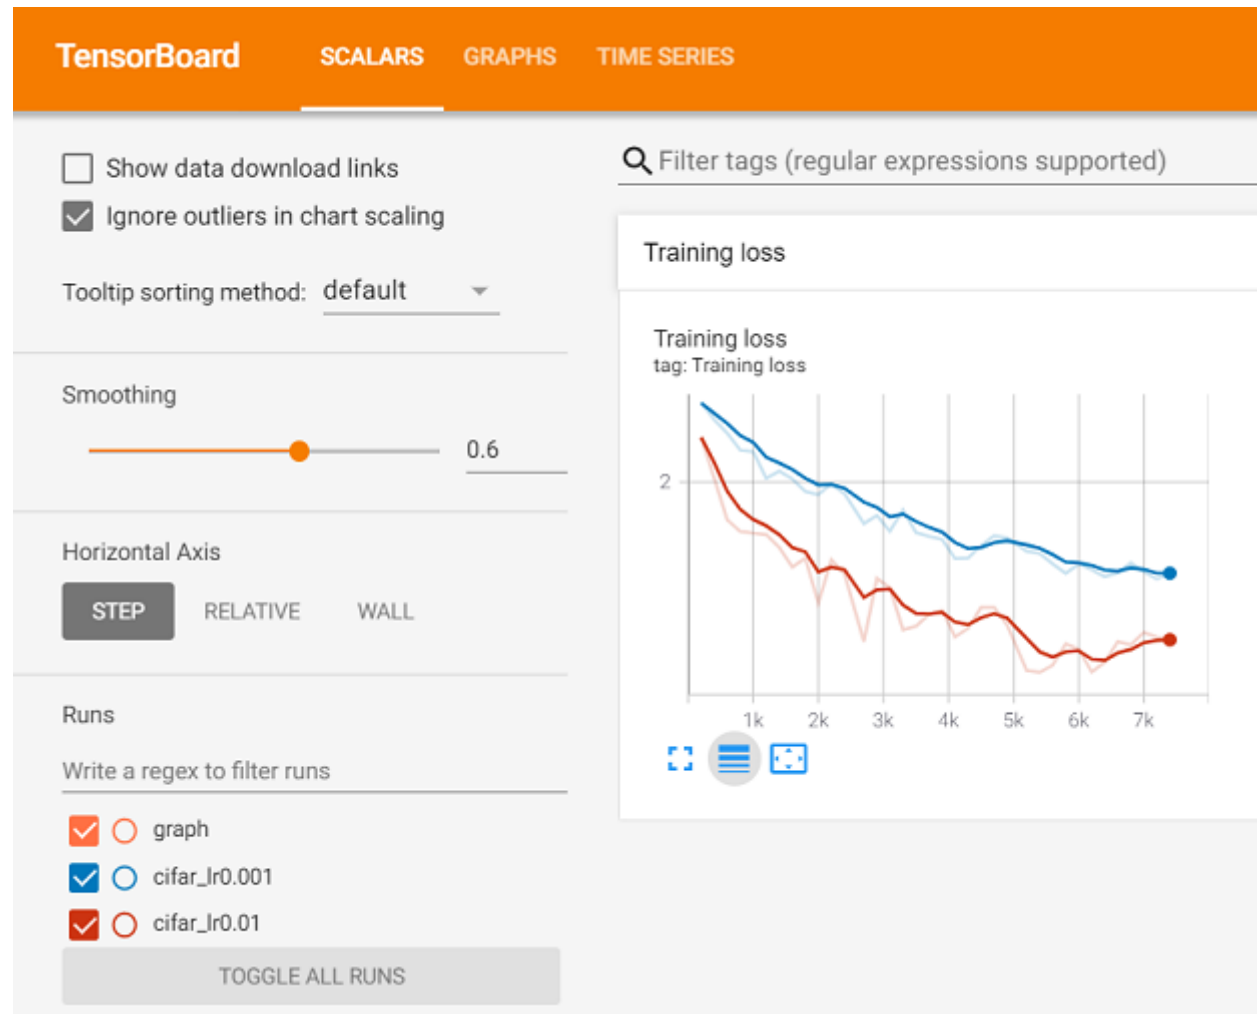
☒ Reload data

Reload Period

30

Pagination Limit

12



5. Grouping Plots

Tensorboard allows you to log multiple information for one experiment. In the following, we shall log 3 different plots:

1. Train loss
2. Test loss

3. Test error

To avoid cluttering the UI, we can **group** the plots by naming them hierarchically. For example, `Loss/train` and `Loss/test` will be grouped together, while `Error/test` will be grouped separately in the TensorBoard interface.

First, let's create our evaluation function to compute the error of a given dataset.

```
In [73]: def evaluate(model, dataloader):
    num_wrong = 0
    test_loss = 0
    total = 0

    criterion = nn.CrossEntropyLoss()

    with torch.no_grad():
        for X, Y in dataloader:
            # transfer input to GPU
            if torch.cuda.is_available():
                X = X.cuda()
                Y = Y.cuda()

            Yhat = model(X)

            # compute test loss
            test_loss += criterion(Yhat, Y).item()

            # compute error
            _, predicted = Yhat.max(axis=1)
            num_wrong += (predicted != Y).sum().item()

            total += len(Y)

    return num_wrong/total, test_loss/total
```

The following code computes the (1) training loss, (2) testing loss and (3) testing error every 100 (batch) iteration. Each time it computes the loss, it will then update tensorboard by calling the command `writer.add_scaler` (lines 56-58) for every values that we want to monitor.

```
In [74]: def train2 (model, writer, lr=0.001, num_epochs=5, loop_per_val=200):

    # set the optimizer
```

```
optimizer = optim.SGD(net.parameters(), lr=lr)

# set the criterion
criterion = nn.CrossEntropyLoss()

# set to training mode
net.train()

# train the model
for e in range(num_epochs):

    running_loss = 0
    running_count = 0

    for i, (inputs, labels) in enumerate(trainloader):

        # transfer input to GPU
        if torch.cuda.is_available():
            inputs = inputs.cuda()
            labels = labels.cuda()

        # set grad to zero
        net.zero_grad()

        # forward propagation
        outputs = net(inputs)

        # compute loss
        loss = criterion(outputs, labels)

        # backward propagation
        loss.backward()

        # update model
        optimizer.step()

        # compute running loss and validation
        running_loss += loss.item()
        running_count += 1

        # display the averaged loss value every 100 cycles
        if i % loop_per_val == loop_per_val - 1:
            train_loss = running_loss / loop_per_val
            running_loss = 0.
```

```

running_count = 0

test_error, test_loss = evaluate(net, testloader)

print(f'[Epoch {e+1:2d}/{num_epochs:d} Iter {e*len(trainloader)+i}]: train_loss = {train_loss:.4f}, test_loss = {t

# Update tensorboard
writer.add_scalar("Group1_Loss/training loss", train_loss, e*len(trainloader)+i)
writer.add_scalar("Group1_Loss/testing loss", test_loss, e*len(trainloader)+i)
writer.add_scalar("Group2_Error/testing error", test_error, e*len(trainloader)+i)

return model

```

Now, let's build the model and log the training loss, test loss and test error.

In [75]:

```

net = build_model()
writer = SummaryWriter('runs/multiscalar_sepgraphs')
net = train2 (net, writer, lr = 0.01, num_epochs=3, loop_per_val=200)
writer.close()

```

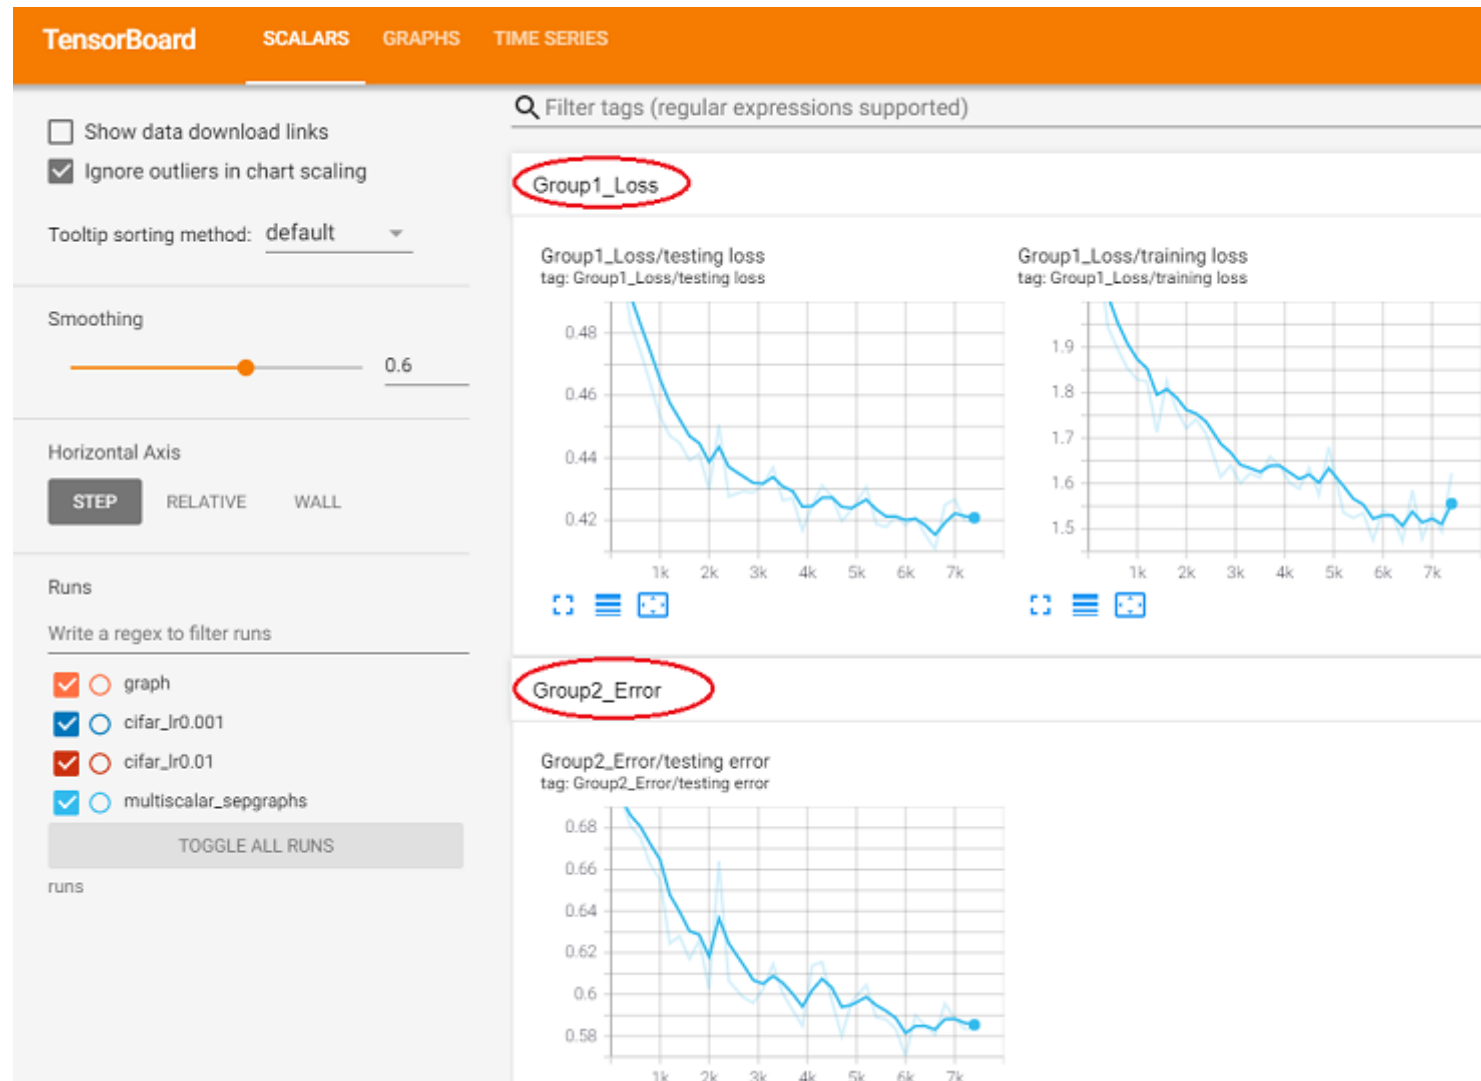
```

[Epoch 1/3 Iter 200/2500]: train_loss = 2.1300, test_loss = 0.5063 | test error: 0.6945
[Epoch 1/3 Iter 400/2500]: train_loss = 1.9450, test_loss = 0.4832 | test error: 0.6810
[Epoch 1/3 Iter 600/2500]: train_loss = 1.8970, test_loss = 0.4743 | test error: 0.6755
[Epoch 1/3 Iter 800/2500]: train_loss = 1.8536, test_loss = 0.4640 | test error: 0.6625
[Epoch 1/3 Iter 1000/2500]: train_loss = 1.8288, test_loss = 0.4534 | test error: 0.6550
[Epoch 1/3 Iter 1200/2500]: train_loss = 1.8248, test_loss = 0.4471 | test error: 0.6245
[Epoch 1/3 Iter 1400/2500]: train_loss = 1.7121, test_loss = 0.4449 | test error: 0.6280
[Epoch 1/3 Iter 1600/2500]: train_loss = 1.8256, test_loss = 0.4392 | test error: 0.6170
[Epoch 1/3 Iter 1800/2500]: train_loss = 1.7620, test_loss = 0.4413 | test error: 0.6260
[Epoch 1/3 Iter 2000/2500]: train_loss = 1.7216, test_loss = 0.4298 | test error: 0.6025
[Epoch 1/3 Iter 2200/2500]: train_loss = 1.7413, test_loss = 0.4507 | test error: 0.6640
[Epoch 1/3 Iter 2400/2500]: train_loss = 1.7100, test_loss = 0.4276 | test error: 0.6065
[Epoch 2/3 Iter 200/2500]: train_loss = 1.6143, test_loss = 0.4294 | test error: 0.5985
[Epoch 2/3 Iter 400/2500]: train_loss = 1.6400, test_loss = 0.4289 | test error: 0.5960
[Epoch 2/3 Iter 600/2500]: train_loss = 1.6002, test_loss = 0.4316 | test error: 0.6025
[Epoch 2/3 Iter 800/2500]: train_loss = 1.6224, test_loss = 0.4369 | test error: 0.6145
[Epoch 2/3 Iter 1000/2500]: train_loss = 1.6128, test_loss = 0.4263 | test error: 0.6005
[Epoch 2/3 Iter 1200/2500]: train_loss = 1.6591, test_loss = 0.4272 | test error: 0.5925
[Epoch 2/3 Iter 1400/2500]: train_loss = 1.6398, test_loss = 0.4168 | test error: 0.5850
[Epoch 2/3 Iter 1600/2500]: train_loss = 1.6019, test_loss = 0.4249 | test error: 0.6140
[Epoch 2/3 Iter 1800/2500]: train_loss = 1.5882, test_loss = 0.4312 | test error: 0.6155
[Epoch 2/3 Iter 2000/2500]: train_loss = 1.6342, test_loss = 0.4275 | test error: 0.5970

```

```
[Epoch 2/3 Iter 2200/2500]: train_loss = 1.5751, test_loss = 0.4198 | test error: 0.5800
[Epoch 2/3 Iter 2400/2500]: train_loss = 1.6807, test_loss = 0.4233 | test error: 0.5965
[Epoch 3/3 Iter 200/2500]: train_loss = 1.5356, test_loss = 0.4308 | test error: 0.6045
[Epoch 3/3 Iter 400/2500]: train_loss = 1.5242, test_loss = 0.4189 | test error: 0.5890
[Epoch 3/3 Iter 600/2500]: train_loss = 1.5350, test_loss = 0.4179 | test error: 0.5880
[Epoch 3/3 Iter 800/2500]: train_loss = 1.4753, test_loss = 0.4212 | test error: 0.5835
[Epoch 3/3 Iter 1000/2500]: train_loss = 1.5407, test_loss = 0.4184 | test error: 0.5705
[Epoch 3/3 Iter 1200/2500]: train_loss = 1.5286, test_loss = 0.4211 | test error: 0.5900
[Epoch 3/3 Iter 1400/2500]: train_loss = 1.4728, test_loss = 0.4154 | test error: 0.5850
[Epoch 3/3 Iter 1600/2500]: train_loss = 1.5848, test_loss = 0.4108 | test error: 0.5805
[Epoch 3/3 Iter 1800/2500]: train_loss = 1.4785, test_loss = 0.4249 | test error: 0.5955
[Epoch 3/3 Iter 2000/2500]: train_loss = 1.5344, test_loss = 0.4268 | test error: 0.5885
[Epoch 3/3 Iter 2200/2500]: train_loss = 1.4938, test_loss = 0.4200 | test error: 0.5835
[Epoch 3/3 Iter 2400/2500]: train_loss = 1.6227, test_loss = 0.4200 | test error: 0.5840
```

Scroll up to the Tensorboard cell to observe the result. You should see something similar to the figure below.



6. Plotting Multiple Plots in the Same Figure

You can also plot multiple scalar data in the same summary. To do that, You need to use the command `writer.addscalars` rather than `writer.addscalar`.


```
In [77]: def train3 (model, writer, lr=0.001, num_epochs=5, loop_per_val=200):

    # set the optimizer
    optimizer = optim.SGD(net.parameters(), lr=lr)

    # set the criterion
    criterion = nn.CrossEntropyLoss()

    # set to training mode
    net.train()

    # train the model
    for e in range(num_epochs):

        running_loss = 0
        running_count = 0

        for i, (inputs, labels) in enumerate(trainloader):

            # transfer input to GPU
            if torch.cuda.is_available():
                inputs = inputs.cuda()
                labels = labels.cuda()

            # set grad to zero
            net.zero_grad()

            # forward propagation
            outputs = net(inputs)

            # compute loss
            loss = criterion(outputs, labels)

            # backward propagation
            loss.backward()

            # update model
            optimizer.step()

            # compute running loss and validation
            running_loss += loss.item()
            running_count += 1
```

```

# display the averaged loss value every 100 cycles
if i % loop_per_val == loop_per_val - 1:
    train_loss = running_loss / loop_per_val
    running_loss = 0.
    running_count = 0

    _, test_loss = evaluate(net, testloader)

    print(f'[Epoch {e+1:2d}/{num_epochs:d} Iter {e*len(trainloader)+i}]: train_loss = {train_loss:.4f}, test_loss = {t

# Update tensorboard
writer.add_scalars('combined', {'train_loss': train_loss,
                                'test_loss': test_loss}, e*len(trainloader)+i)

return model

```

Now, let's build the model and log the training loss and test loss.

In [79]:

```

net = build_model()
writer = SummaryWriter('runs/multiscalar_samegraph')
net = train3 (net, writer, lr = 0.01, num_epochs=3, loop_per_val=200)
writer.close()

```

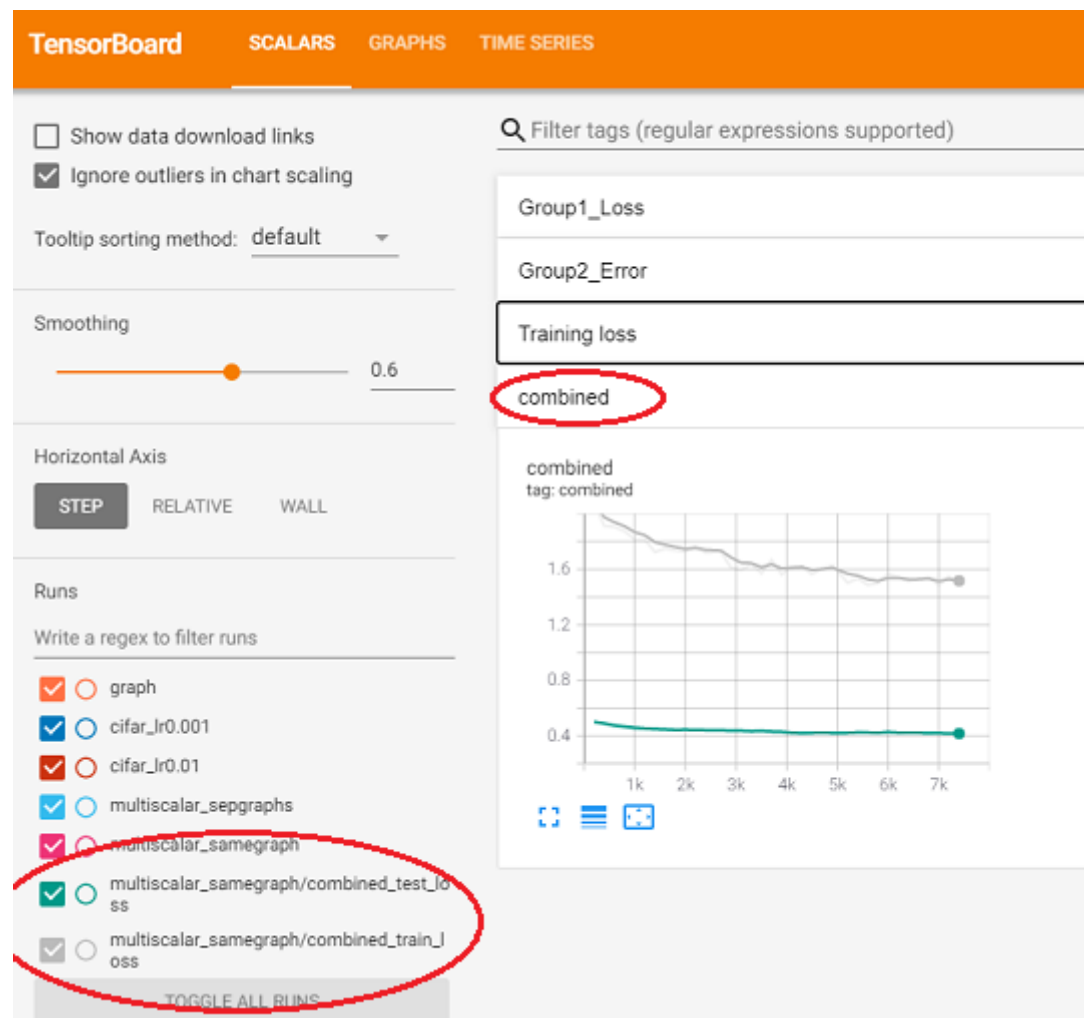
```

[Epoch 1/3 Iter 200/2500]: train_loss = 2.1144, test_loss = 0.5067
[Epoch 1/3 Iter 400/2500]: train_loss = 1.9488, test_loss = 0.4768
[Epoch 1/3 Iter 600/2500]: train_loss = 1.8901, test_loss = 0.4655
[Epoch 1/3 Iter 800/2500]: train_loss = 1.7990, test_loss = 0.4591
[Epoch 1/3 Iter 1000/2500]: train_loss = 1.7821, test_loss = 0.4588
[Epoch 1/3 Iter 1200/2500]: train_loss = 1.8275, test_loss = 0.4530
[Epoch 1/3 Iter 1400/2500]: train_loss = 1.7208, test_loss = 0.4498
[Epoch 1/3 Iter 1600/2500]: train_loss = 1.7355, test_loss = 0.4377
[Epoch 1/3 Iter 1800/2500]: train_loss = 1.7926, test_loss = 0.4365
[Epoch 1/3 Iter 2000/2500]: train_loss = 1.7570, test_loss = 0.4312
[Epoch 1/3 Iter 2200/2500]: train_loss = 1.7719, test_loss = 0.4440
[Epoch 1/3 Iter 2400/2500]: train_loss = 1.7608, test_loss = 0.4365
[Epoch 2/3 Iter 200/2500]: train_loss = 1.6473, test_loss = 0.4352
[Epoch 2/3 Iter 400/2500]: train_loss = 1.6725, test_loss = 0.4281
[Epoch 2/3 Iter 600/2500]: train_loss = 1.5417, test_loss = 0.4406
[Epoch 2/3 Iter 800/2500]: train_loss = 1.6487, test_loss = 0.4380
[Epoch 2/3 Iter 1000/2500]: train_loss = 1.6937, test_loss = 0.4336
[Epoch 2/3 Iter 1200/2500]: train_loss = 1.6178, test_loss = 0.4268
[Epoch 2/3 Iter 1400/2500]: train_loss = 1.5901, test_loss = 0.4260

```

```
[Epoch 2/3 Iter 1600/2500]: train_loss = 1.6707, test_loss = 0.4315
[Epoch 2/3 Iter 1800/2500]: train_loss = 1.6844, test_loss = 0.4273
[Epoch 2/3 Iter 2000/2500]: train_loss = 1.5781, test_loss = 0.4281
[Epoch 2/3 Iter 2200/2500]: train_loss = 1.5894, test_loss = 0.4180
[Epoch 2/3 Iter 2400/2500]: train_loss = 1.6271, test_loss = 0.4222
[Epoch 3/3 Iter 200/2500]: train_loss = 1.5054, test_loss = 0.4287
[Epoch 3/3 Iter 400/2500]: train_loss = 1.4911, test_loss = 0.4176
[Epoch 3/3 Iter 600/2500]: train_loss = 1.5010, test_loss = 0.4134
[Epoch 3/3 Iter 800/2500]: train_loss = 1.5056, test_loss = 0.4175
[Epoch 3/3 Iter 1000/2500]: train_loss = 1.5114, test_loss = 0.4230
[Epoch 3/3 Iter 1200/2500]: train_loss = 1.5527, test_loss = 0.4416
[Epoch 3/3 Iter 1400/2500]: train_loss = 1.5537, test_loss = 0.4209
[Epoch 3/3 Iter 1600/2500]: train_loss = 1.5315, test_loss = 0.4139
[Epoch 3/3 Iter 1800/2500]: train_loss = 1.5627, test_loss = 0.4330
[Epoch 3/3 Iter 2000/2500]: train_loss = 1.5188, test_loss = 0.4167
[Epoch 3/3 Iter 2200/2500]: train_loss = 1.5444, test_loss = 0.4262
[Epoch 3/3 Iter 2400/2500]: train_loss = 1.5945, test_loss = 0.4136
```

Scroll up to the Tensorboard cell to observe the result. You should see something similar to the figure below.



8. Other uses of Tensorboard

Besides, you can explore the following output to tensorboard:

1. Displaying images via the `writer.add_image` method
2. Displaying audio via the `writer.add_audio` method

3. Displaying video via the `writer.add_video` method
4. Displaying histogram via the `writer.add_histogram` method
5. Disualize the lower dimensional representation of higher dimensional data through the `add_embedding` method